

抢渡长江的数学建模分析

李俊霖

2025 年 3 月 14 日

1 任务


我将用 python 构建一个基于 llama 2 的聊天机器人，然后用 Api 对于它进行调用

2 结构

用 Streamlit 作为前端，用 Python 构建一个 Llama 2 聊天机器人，而 LLM 后端则通过对 Replicate 上托管的 Llama 2 模型的 API 调用来处理

3 获取 Replicate API 令牌

- (1) 转到 Replicate <https://llama2.streamlit.app>
- (2) 使用您的 GitHub 帐户登录。
- (3) 转到 API 令牌页面并复制您的 API 令牌。



新建一个 Api-Token，这个 Api-Token 就是之后你进行远程调用时所需要的

4 设置编码环境

本地开发

要设置本地编码环境，请在命令行提示符中输入以下命令：

```
1 pip install streamlit replicate
```

5 构建应用程序

总代码

```
1 import streamlit as st
2 import replicate
3 import os
4
5 # App title
6 st.set_page_config(page_title=" Llama_2_Chatbot")
7
8 # Replicate Credentials
9 with st.sidebar:
10     st.title(' Llama_2_Chatbot')
11     if 'REPLICATE_API_TOKEN' in st.secrets:
12         st.success('API_key_already_provided!', icon=' ')
13         replicate_api = st.secrets['REPLICATE_API_TOKEN']
14     else:
15         replicate_api = st.text_input('Enter_Replicate_API_
16                                     token:', type='password')
17         if not (replicate_api.startswith('r8_') and len(
18             replicate_api)==40):
19             st.warning('Please_enter_your_credentials!', icon='
20             ')
21         else:
22             st.success('Proceed_to_entering_your_prompt_message!
23             ', icon=' ')
```

```

20 os.environ['REPLICATE_API_TOKEN'] = replicate_api
21
22 st.subheader('Models and parameters')
23 selected_model = st.sidebar.selectbox('Choose a Llama2
    model', ['Llama2-7B', 'Llama2-13B'], key='
    selected_model')
24 if selected_model == 'Llama2-7B':
25     llm = 'a16z-infra/llama7b-v2-chat:4
        f0a4744c7295c024a1de15e1a63c880d3da035fa1f49bfd344fe076074c8eea
        '
26 elif selected_model == 'Llama2-13B':
27     llm = 'a16z-infra/llama13b-v2-chat:
        df7690f1994d94e96ad9d568eac121aecf50684a0b0963b25a41cc40061269e5
        '
28 temperature = st.sidebar.slider('temperature', min_value
    =0.01, max_value=5.0, value=0.1, step=0.01)
29 top_p = st.sidebar.slider('top_p', min_value=0.01,
    max_value=1.0, value=0.9, step=0.01)
30 max_length = st.sidebar.slider('max_length', min_value=32,
    max_value=128, value=120, step=8)
31 st.markdown(' Learn how to build this app in this [blog](
    https://blog.streamlit.io/how-to-build-a-llama-2-
    chatbot/)!')
32
33 # Store LLM generated responses
34 if "messages" not in st.session_state.keys():
35     st.session_state.messages = [{"role": "assistant", "
        content": "How may I assist you today?"}]
36
37 # Display or clear chat messages
38 for message in st.session_state.messages:
39     with st.chat_message(message["role"]):
40         st.write(message["content"])

```

```

41
42 def clear_chat_history():
43     st.session_state.messages = [{"role": "assistant", "
        content": "How may I assist you today?"}]
44 st.sidebar.button('Clear Chat History', on_click=
        clear_chat_history)
45
46 # Function for generating LLaMA2 response. Refactored from
        https://github.com/a16z-infra/llama2-chatbot
47 def generate_llama2_response(prompt_input):
48     string_dialogue = "You are a helpful assistant. You do not
        respond as 'User' or pretend to be 'User'. You only
        respond once as 'Assistant'."
49     for dict_message in st.session_state.messages:
50         if dict_message["role"] == "user":
51             string_dialogue += "User: " + dict_message["content"]
                + "\n\n"
52         else:
53             string_dialogue += "Assistant: " + dict_message["
                content"] + "\n\n"
54     output = replicate.run('a16z-infra/llama13b-v2-chat:
        df7690f1994d94e96ad9d568eac121aecf50684a0b0963b25a41cc40061269e5
        ',
55                             input={"prompt": f"{string_dialogue}{
                prompt_input}Assistant:"},
56                             "temperature":temperature, "top_p
                ":top_p, "max_length":
                max_length, "
                repetition_penalty":1})
57     return output
58
59 # User-provided prompt
60 if prompt := st.chat_input(disabled=not replicate_api):

```

```

61     st.session_state.messages.append({"role": "user", "content
        ": prompt})
62     with st.chat_message("user"):
63         st.write(prompt)
64
65     # Generate a new response if last message is not from
        assistant
66     if st.session_state.messages[-1]["role"] != "assistant":
67         with st.chat_message("assistant"):
68             with st.spinner("Thinking..."):
69                 response = generate_llama2_response(prompt)
70                 placeholder = st.empty()
71                 full_response = ''
72                 for item in response:
73                     full_response += item
74                     placeholder.markdown(full_response)
75                 placeholder.markdown(full_response)
76     message = {"role": "assistant", "content": full_response}
77     st.session_state.messages.append(message)

```


6 我部署时解决的问题

- (1) 首先，创建一个文件夹，用 vscode 打开文件
- (2) 设置.env 虚拟环境
- (3) 然后按照部署相关环境，比如 streamlit 和 republic
- (4) 之后在此文件夹中新建一个文件夹.streamlit
- (5) 在.streamlit 文件夹中新建 secrets.toml 文件，然后如照片那样在其中输入之前的 Api
- (6) 最后进行调用，记住一定要用

```

1     streamlit run main.py

```



最后成功了

7 声明

此篇文章深度借鉴了 <https://blog.streamlit.io/how-to-build-a-llama-2-chatbot/> 如何构建 Llama 2 聊天机器人这篇文章，这篇文章只是记录一下自己的学习过程与学习心得，希望对于有所帮助